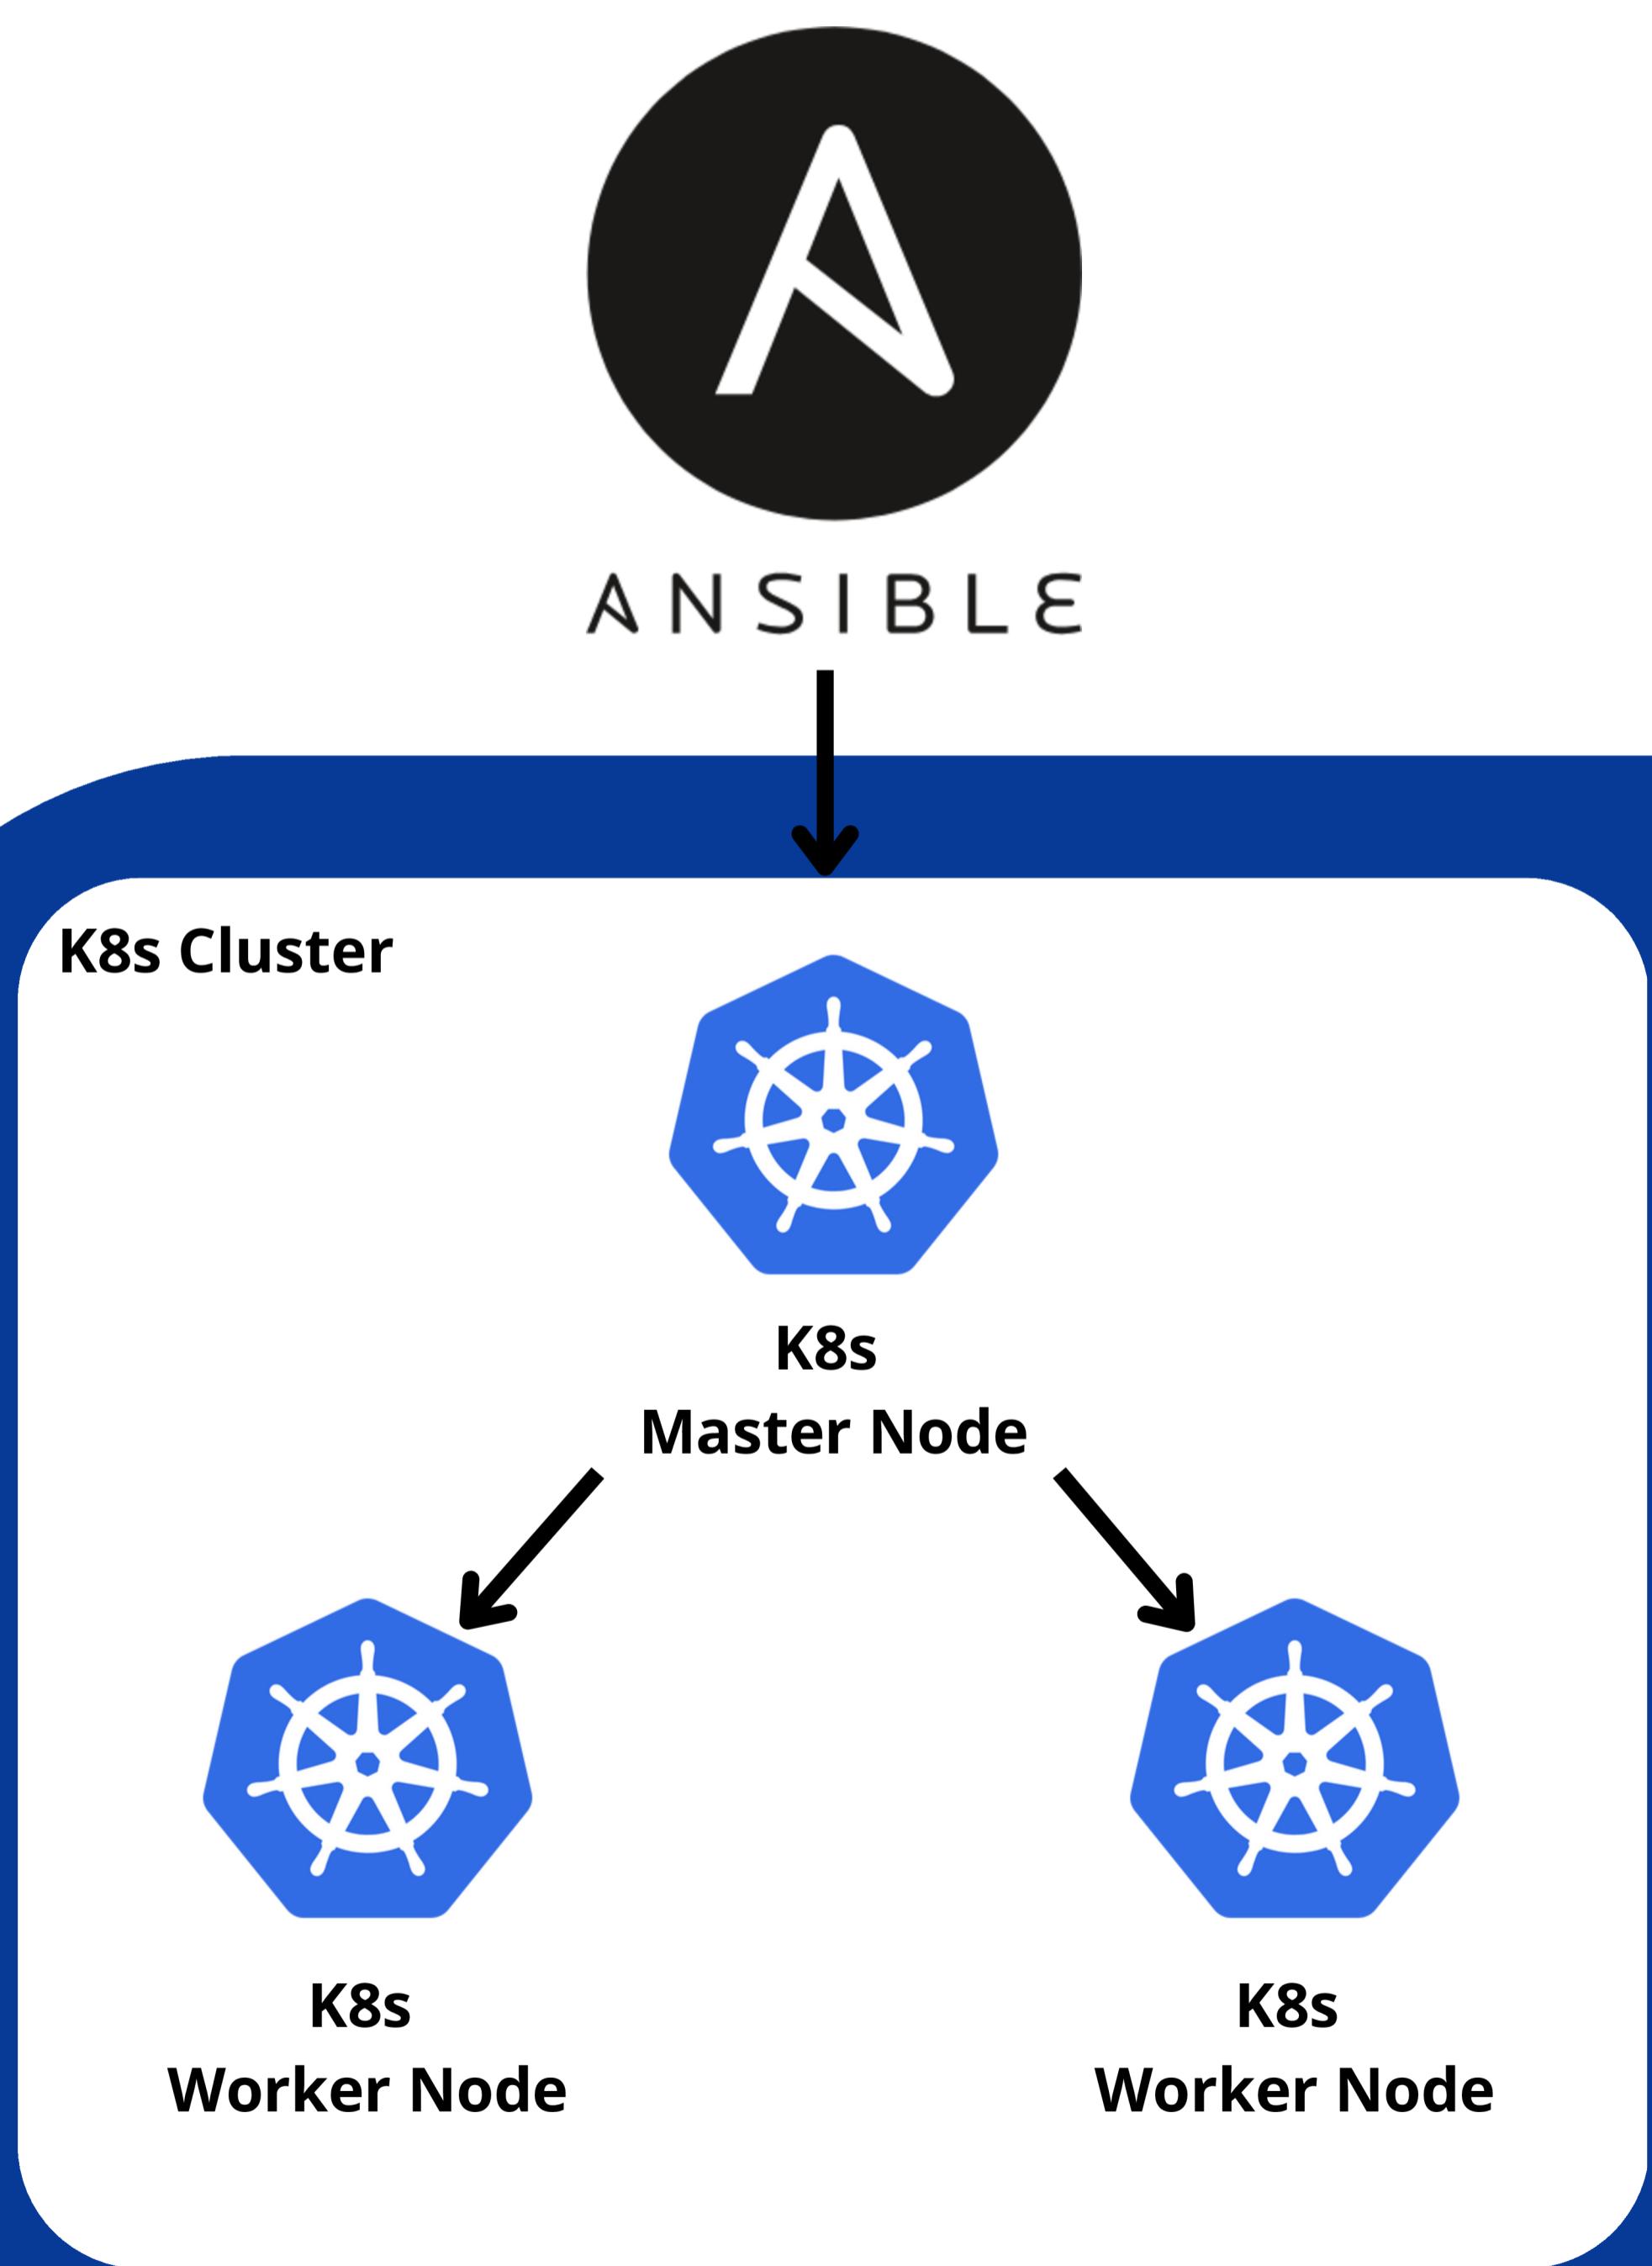


Ansible for Automating Kubernetes Cluster Installation



Introduction to Ansible:



Ansible is an open-source automation tool that simplifies IT automation, configuration management, application deployment, and task execution. It allows administrators to write simple scripts, known as playbooks, which can manage and configure multiple systems simultaneously. Ansible's agentless architecture makes it particularly appealing for environments where installing management agents is impractical.

Advantages:



- **Agentless Architecture:** Ansible does not require any software to be installed on the target machines, making it easier to maintain.
- **Declarative Language:** Playbooks are written in YAML, a human-readable language that allows IT professionals to define the desired state of systems.
- **Idempotency:** Ansible ensures that running a playbook multiple times results in the same state, preventing unintended changes.

Why Ansible?



Ansible offers a simple, agentless approach to automation. It allows for idempotent operations, ensuring that tasks are executed only if necessary, which is crucial for maintaining the desired state of your infrastructure.

Setup:



Inventory Management

In Ansible, an inventory file is used to define the hosts and groups of hosts on which tasks will be executed. For our Kubernetes cluster, we will have a master node and multiple worker nodes, all of which need to be listed in the inventory.

Example “inventory.yml” file :

```
[masters]
master1 ansible_host=192.168.1.10 ansible_user=ubuntu

[workers]
worker1 ansible_host=192.168.1.20 ansible_user=ubuntu
worker2 ansible_host=192.168.1.21 ansible_user=ubuntu
```

Explanation:

- **[masters]:** This group contains all master nodes. In this example, there is one master node (master1) with the IP address 192.168.1.10. The ansible_user is set to ubuntu, which is the user that Ansible will use to SSH into the master node.
- **[workers]:** This group contains all worker nodes. Here, we have two worker nodes (worker1 and worker2) with IP addresses 192.168.1.20 and 192.168.1.21, respectively. The ansible_user is also set to ubuntu for these nodes.



Playbook Overview:

Structure of a Playbook

A playbook is a YAML file that defines a series of tasks to be executed on a group of hosts. It is composed of plays, each targeting a group of hosts, and contains tasks that will be performed on those hosts.

Key Components:

- **Hosts:** The target machines where tasks will be executed.
- **Tasks:** Individual actions to be performed, such as installing software, configuring files, or starting services.
- **Handlers:** Tasks triggered by other tasks, typically used for restarting services.
- **Variables:** Parameters that can be reused across tasks, making playbooks more flexible.

Example Playbook: Installing and Starting Nginx :

```
- name: Install and start Nginx on all nodes
  hosts: all
  become: yes

  tasks:
    - name: Update apt cache
      apt:
        update_cache: yes

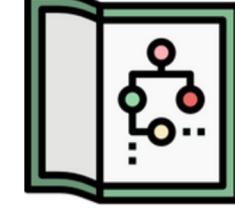
    - name: Install Nginx
      apt:
        name: nginx
        state: present

    - name: Ensure Nginx is running and enabled at boot
      service:
        name: nginx
        state: started
        enabled: yes
```

Explanation:

- **hosts: all:** The playbook will run on all hosts defined in the inventory file, including both master and worker nodes.
- **become: yes:** The tasks will be executed with elevated (root) privileges.
- **Tasks:**
 - **Update apt cache:** This task ensures that the package list is up to date.
 - **Install Nginx:** Installs the Nginx web server if it is not already present.
 - **Ensure Nginx is running and enabled at boot:** Starts the Nginx service and ensures that it will automatically start on system boot.

This example demonstrates the basic structure of an Ansible playbook. It shows how to define tasks, use Ansible modules like apt and service, and target multiple hosts. The simplicity of this example makes it a good introduction before diving into more complex playbooks for Kubernetes cluster setup.



Kubernetes Installation Playbook:

This section provides a step-by-step explanation of the single consolidated playbook used to automate the installation of a Kubernetes cluster. This playbook combines all necessary tasks, from preparing the system to initializing the cluster.

"Kubernetes.yml" playbook :

```
- name: Set up Kubernetes Cluster
  hosts: all
  become: yes

  tasks:
    - name: Disable swap
      command: swapoff -a

    - name: Remove swap entry from /etc/fstab
      replace:
        path: /etc/fstab
        regexp: '^(.*\sswap\s.*$)'
        replace: '#\1'

    - name: Install required packages
      apt:
        name:
          - apt-transport-https
          - ca-certificates
          - curl
          - gnupg
          - lsb-release
        state: present
        update_cache: yes

    - name: Add Kubernetes GPG key
      apt_key:
        url: https://packages.cloud.google.com/apt/doc/apt-key.gpg
        state: present

    - name: Add Kubernetes apt repository
      apt_repository:
        repo: deb [signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg] https://apt.kubernetes.io/ kubernetes-xenial main
        state: present
        filename: kubernetes

    - name: Install Kubernetes components
      apt:
        name:
          - kubelet
          - kubeadm
          - kubectl
        state: present
        update_cache: yes
```

```
- name: Hold Kubernetes packages at the current version
  apt:
    name:
      - kubelet
      - kubeadm
      - kubectl
    state: present
    lock: yes

- name: Install containerd
  apt:
    name: containerd
    state: present

- name: Create containerd configuration directory
  file:
    path: /etc/containerd
    state: directory

- name: Generate default containerd configuration
  command: containerd config default > /etc/containerd/config.toml

- name: Configure containerd to use systemd as cgroup driver
  replace:
    path: /etc/containerd/config.toml
    regexp: 'SystemdCgroup = false'
    replace: 'SystemdCgroup = true'

- name: Restart containerd
  service:
    name: containerd
    state: restarted
    enabled: yes

- name: Initialize Kubernetes master node
  hosts: masters
  become: yes

  tasks:
    - name: Initialize Kubernetes master node
      command: kubeadm init --pod-network-cidr=192.168.0.0/16
      register: init_output
```

```

- name: Create .kube directory for root
  file:
    path: /root/.kube
    state: directory
    mode: '0755'

- name: Copy kubeconfig to root's .kube directory
  command: "{{ item }}"
  with_items:
    - mkdir -p $HOME/.kube
    - cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
    - chown $(id -u):$(id -g) $HOME/.kube/config

- name: Save the join command to a file
  copy:
    content: "{{ init_output.stdout }}"
    dest: /root/join_command.txt

- name: Join worker nodes to the cluster
  hosts: workers
  become: yes

  tasks:
    - name: Fetch join command from master
      command: cat /root/join_command.txt
      delegate_to: master1
      register: join_command

    - name: Join worker node to the Kubernetes cluster
      command: "{{ join_command.stdout_lines }}"

```

This playbook is designed to automate the entire process of setting up a Kubernetes cluster. It handles the configuration and installation steps required for both master and worker nodes, ensuring a consistent and reliable deployment. The playbook is divided into three main sections: Global Setup on All Nodes, Master Node Setup, and Worker Node Setup.

1. Global Setup on All Nodes:

- Disable Swap:** Ensures that swap is turned off on all nodes, which is a requirement for Kubernetes.
- Remove Swap Entry:** Comments out any swap entries in /etc/fstab to prevent swap from being re-enabled on reboot.
- Install Required Packages:** Installs essential packages like apt-transport-https, curl, and others needed to add repositories and install Kubernetes components.
- Add Kubernetes GPG Key and Repository:** Adds the GPG key and the Kubernetes APT repository to ensure secure and reliable installation of Kubernetes packages.
- Install Kubernetes Components:** Installs kubelet, kubeadm, and kubectl on all nodes, which are the core components required to set up a Kubernetes cluster.
- Hold Kubernetes Package Versions:** Locks the installed Kubernetes packages at their current versions to prevent unintended upgrades.
- Install and Configure Containerd:** Installs the containerd runtime, creates its configuration directory, and sets it up to use systemd as the cgroup driver. Containerd is then restarted to apply the changes.

2. Master Node Setup:

- **Initialize Kubernetes Master Node:** Initializes the Kubernetes control plane on the master node using `kubeadm`, specifying a pod network CIDR.
- **Configure Kubeconfig for Root:** Sets up the `kubeconfig` file in the root user's `.kube` directory to allow command-line interaction with the Kubernetes cluster.
- **Save Join Command:** Saves the command needed for worker nodes to join the cluster into a file for later use.

3. Worker Node Setup:

- **Fetch Join Command:** Retrieves the join command saved on the master node to be used on the worker nodes.
- **Join Worker Nodes:** Executes the join command on each worker node to integrate them into the Kubernetes cluster.

How I execute the Playbook ?

To execute the consolidated playbook, follow these steps:

1. Ensure SSH Access: Make sure you can SSH into all the nodes listed in your inventory.
2. Run the Playbook: Execute the playbook using the following command:



`ansible-playbook -i inventory.yml kubernetes.yml`



3. Verify the Installation: After the playbook completes, verify that Kubernetes is installed correctly by running:

`kubectl get nodes`



Here is an example of the output from the "kubectl get nodes" command, which provides information about the nodes in the Kubernetes cluster:

```
master01@master-node:~$ kubectl get nodes
NAME           STATUS    ROLES      AGE     VERSION
master-node    Ready     control-plane   42h    v1.30.2
worker-node    Ready     <none>        25h    v1.30.2
worker-node02  Ready     <none>        86s    v1.30.2
```

Best Practices:

Here are some essential best practices to use in Ansible:



- **Use Roles:** Organize tasks into roles to promote reusability and maintainability.
- **Modular Playbooks:** Break down large playbooks into smaller, focused ones for better readability and management.
- **Idempotency:** Ensure tasks are idempotent, meaning they can be run multiple times without causing unintended changes.
- **Use Variables:** Centralize and manage configuration with variables to make playbooks more flexible.
- **Secure Sensitive Data:** Use Ansible Vault to encrypt sensitive information such as passwords and keys.
- **Version Control:** Track changes to playbooks and roles using a version control system like Git.
- **Test Playbooks:** Regularly test playbooks in a staging environment before applying them to production.
- **Consistent Naming Conventions:** Follow consistent naming conventions for tasks, variables, and roles to improve clarity and reduce errors.
- **Documentation:** Document playbooks and roles to ensure they are understandable and maintainable by others.
- **Error Handling:** Implement proper error handling to manage failures gracefully during execution.

Conclusion:

In conclusion, this document provides a streamlined approach to automating Kubernetes cluster setup using Ansible. By consolidating the setup process into a unified playbook, it offers a clear and efficient pathway to deploy and manage Kubernetes clusters with best practices in mind.

The playbook and its explanations aim to simplify complex tasks, ensuring a robust and reliable Kubernetes environment. Implementing these practices can enhance consistency, security, and efficiency in your infrastructure management.

I hope you find this guide valuable for your own Kubernetes projects and automation needs. Thank you for taking the time to review it, and feel free to reach out with any questions or feedback.



Omar
MASMOUDI

Thank
you!